

# Lecture 9-10

## Object Categorisation

### Bag of Words

### Maximum Margin Classifier

Tae-Kyun Kim

# Object Categorisation - Challenges



Illumination



Poses/orientations



Clutter



Occlusions



intra-class variations



View-points


# Image Representation

## - Bag of Words



A good image representation is required.

Image are often represented as finite dimensional vectors.

Later, the vectors are classified into object categories.

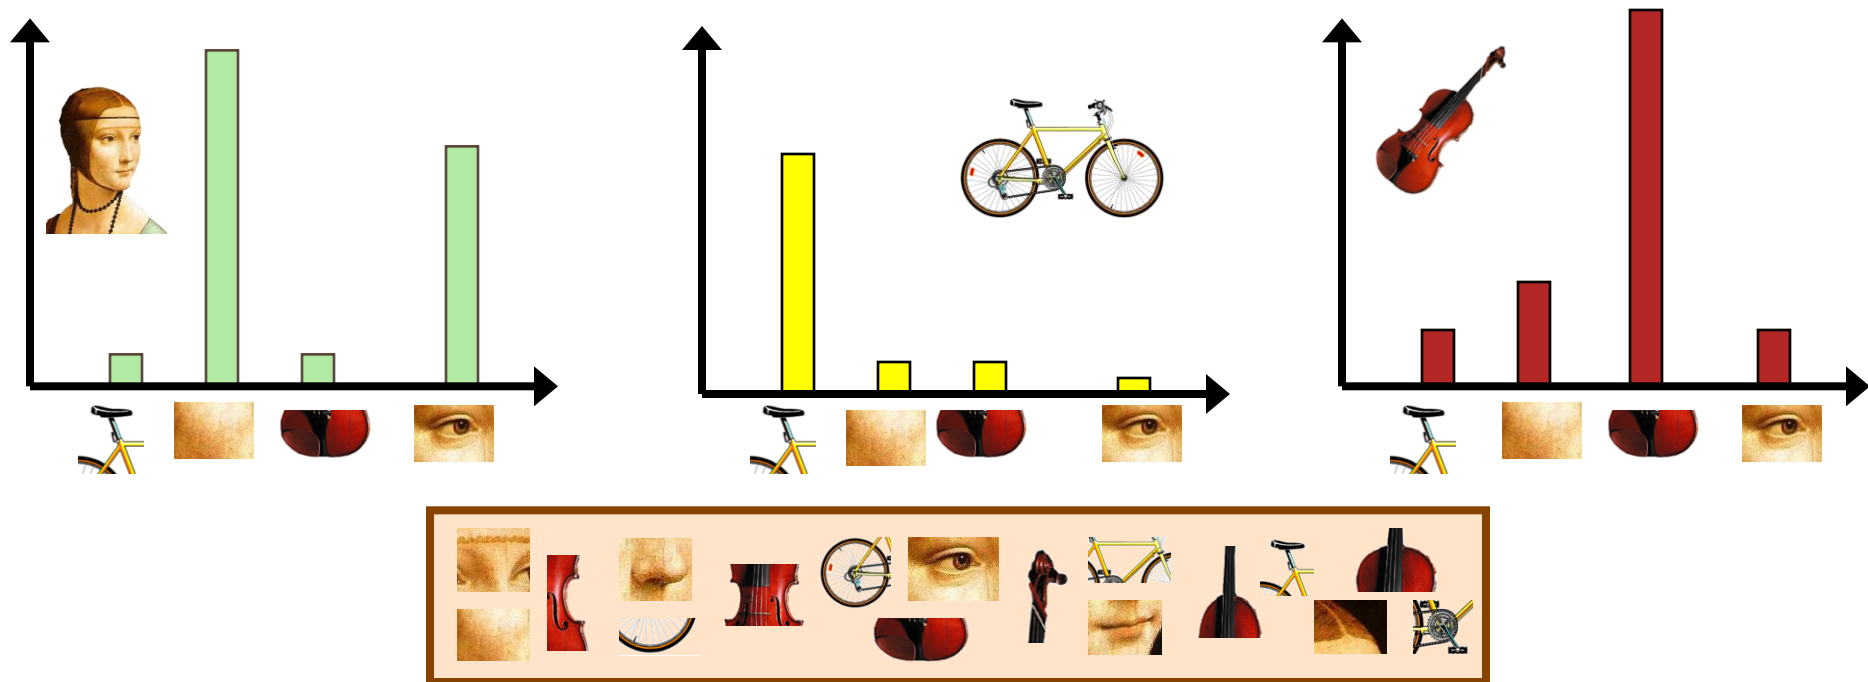


# Bag of Words Model



# Bag of Words

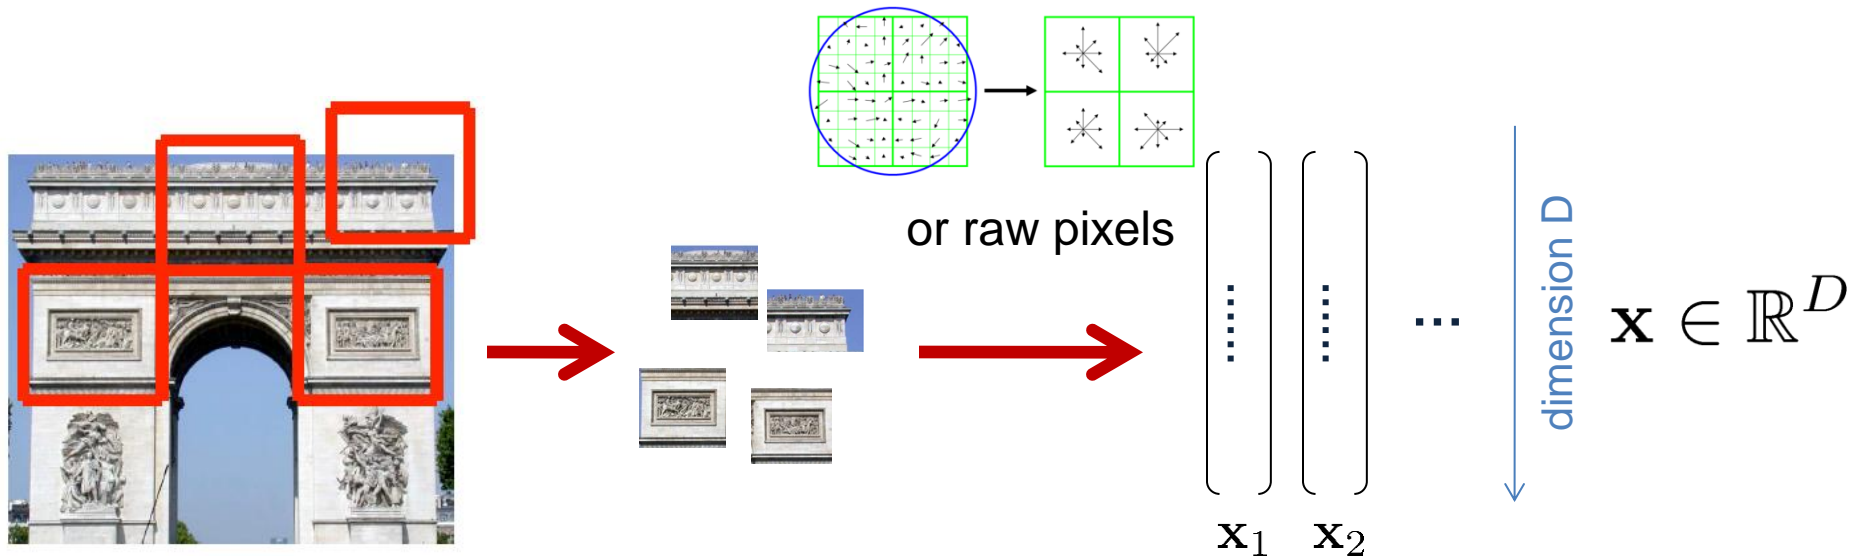
- Independent features (code words)
- Histogram representation






# Visual Words

- Visual words are base elements to describe an image.
- Interest points are detected from an image.
  - E.g. Corners, Blob detector, SIFT (Scale-Invariant Feature Transform) detector (<http://www.cs.ubc.ca/~lowe/keypoints/>)
- Image patches are collected and represented by descriptors.
  - SIFT or raw pixel intensity



# Visual Words

- Instead of a sparse Interest point detector, we can use a *dense grid*.
- Image patches are collected around *all points* on the grid.




# Building a Visual Codebook (Dictionary)

- Visual words (real-valued vectors) can be compared using Euclidean distance:

$$E(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{\sum_d (\mathbf{x}_1^d - \mathbf{x}_2^d)^2}$$

- These vectors are divided into groups which are similar, essentially clustered together, to form a codebook.






# K-means Clustering for Codebook


Lecture 3-4

- A large number of image patches are collected from a representative set of images.
- The two steps repeat until no change in membership:
  1. Compute the center of each cluster as the data mean of the respective cluster
  2. Reassign each data point to the cluster whose center is nearest
- The cluster centers (mean vectors) form a visual dictionary.




# Histogram of Visual Words


- Every visual word is compared with codewords and assigned to the nearest codeword.
- Histogram bins are codewords and each bin counts the number of words assigned to the codeword.



Nearest  
Neighbour  
Matching



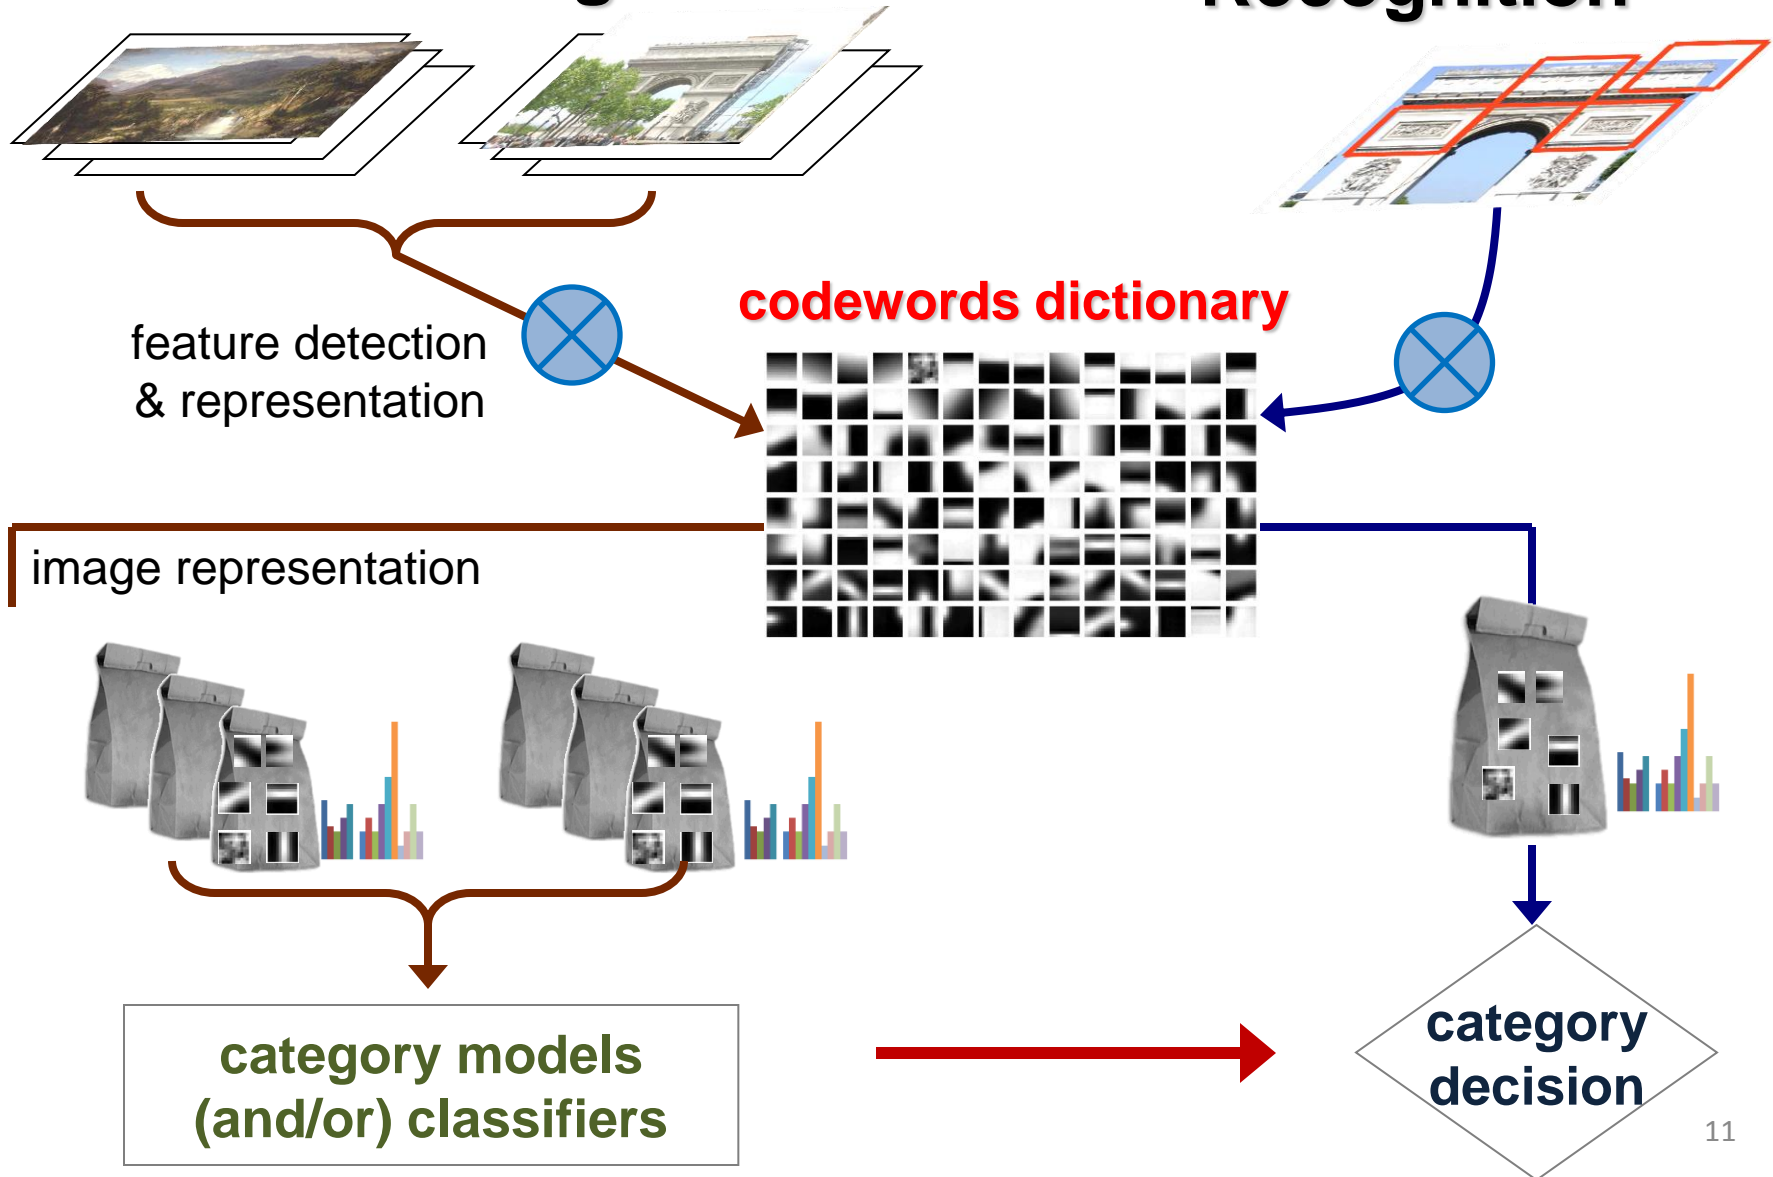
**“bag of words”**



# Categorisation

## Learning

## Recognition



# Summary of K-means Codebook

- The histogram is K-D, a highly compact and robust representation of images.
- The histogram representation greatly facilitates modelling images and categorising them.
- Codeword assignment (quantisation process) by K-means is time-demanding.
- K-means is an unsupervised learning method.

# Sparse Kernel Machine (Maximum Margin Classifier or Support Vector Machine)

$$f : \mathbf{x} \in \mathbb{R}^D \rightarrow t \in \{1, \dots, n\}$$



# Linear Discriminant Functions

- A discriminant function maps an input vector  $\mathbf{x}$  into one of  $K$  classes, denoted  $C_k$ . For simplicity, consider two classes.
- *Linear* discriminant function takes the form of

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

where  $\mathbf{w}$  is a *weight vector* and  $w_0$  a *bias*.

- A vector  $\mathbf{x}$  is assigned to  $C_1$  if  $y(\mathbf{x}) \geq 0$ , and  $C_2$  otherwise.

- The *decision boundary* is defined by  $y(\mathbf{x}) = 0$ , which is a  $(D-1)$ -dimensional hyperplane in the  $D$ -dimensional input space.
- Consider  $\mathbf{x}_A$  and  $\mathbf{x}_B$  on the decision surface.

→  $y(\mathbf{x}_A) = y(\mathbf{x}_B) = 0$

→  $\mathbf{w}^T(\mathbf{x}_A - \mathbf{x}_B) = 0$

The vector  $\mathbf{w}$  is orthogonal to the decision surface.  $\mathbf{w}$  determines the *direction* of the decision surface.

- For  $\mathbf{x}$  on the decision surface,  $y(\mathbf{x}) = 0$ , so the normal distance from the origin to the decision surface is

$$\frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\|} = -\frac{w_0}{\|\mathbf{w}\|}$$


Thus,  $w_0$  determines the *location* of the decision surface.

- For an arbitrary point  $\mathbf{x}$ , its orthogonal projection onto the decision surface  $\mathbf{x}_\perp$  is such that

$$\mathbf{x} = \mathbf{x}_\perp + r \frac{\mathbf{w}}{\|\mathbf{w}\|}.$$


- Multiply both sides by  $\mathbf{w}^\top$  and add  $w_0$ , and use  $y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$  and  $y(\mathbf{x}_\perp) = \mathbf{w}^\top \mathbf{x}_\perp + w_0 = 0$ , we have

$$r = \frac{y(\mathbf{x})}{\|\mathbf{w}\|}.$$



# Kernel Trick

- We often need a *nonlinear* decision boundary.
- The input space is transformed into a *high-dimensional feature space* by a nonlinear mapping  $\phi : \mathbf{x} \rightarrow \phi(\mathbf{x})$ , where a linear decision boundary can separate all data points.





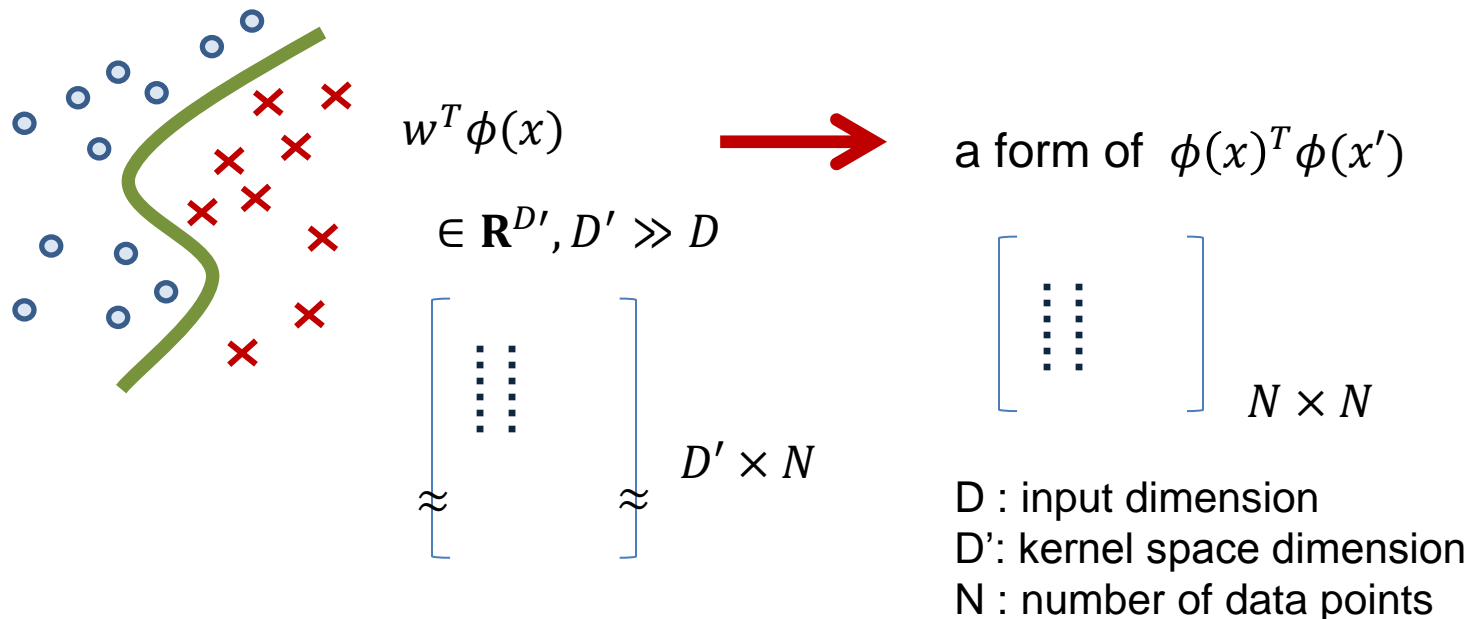
## *Kernel trick:*

The concept of a kernel formulated as an inner product in a feature space allows to build nonlinear extensions of many ML algorithms.

- Many ML problems can be recast into an equivalent *dual representation*, where the predictions are based on linear combinations of a *kernel function* evaluated at training data points.

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

- The kernel is a symmetric function  $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$



- An example is

$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z})^2.$$

If  $\mathbf{x} = (x_1, x_2)$ ,  $\mathbf{z} = (z_1, z_2)$ ,

$$\begin{aligned} k(\mathbf{x}, \mathbf{z}) &= (\mathbf{x}^T \mathbf{z})^2 = (x_1 z_1 + x_2 z_2)^2 \\ &= x_1^2 z_1^2 + 2x_1 z_1 x_2 z_2 + x_2^2 z_2^2 \\ &= \begin{pmatrix} x_1^2, \sqrt{2}x_1 x_2, x_2^2 \end{pmatrix} \begin{pmatrix} z_1^2, \sqrt{2}z_1 z_2, z_2^2 \end{pmatrix}^T \\ &= \phi(\mathbf{x})^T \phi(\mathbf{z}). \end{aligned}$$

The feature mapping takes the form  $\phi(\mathbf{x}) = (x_1^2, \sqrt{2}x_1 x_2, x_2^2)^T$  comprising all second order terms.

- A necessary and sufficient condition for a valid kernel  $k(\mathbf{x}, \mathbf{x}')$ :

$$\mathbf{x}^T \mathbf{K} \mathbf{x} \geq 0$$

where  $\mathbf{K}$  is the Gram matrix whose elements are  $k(\mathbf{x}_i, \mathbf{x}_j')$  i.e.  $\in \mathbf{R}^{N \times N}$  ( $N$  is the number of training data points).  $\mathbf{K}$  should be positive semidefinite.

- *Typical examples of Kernels:*

Linear kernel:  $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$

Polynomial kernel:  $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + c)^M$  with  $c > 0$

Gaussian kernel:  $k(\mathbf{x}, \mathbf{x}') = \exp(-||\mathbf{x} - \mathbf{x}'||^2 / 2\sigma^2)$

# Maximum Margin Classifiers

- *Sparse Kernel Machine*: has sparse solutions i.e. the predictions for new inputs depend only on the kernel function evaluated at a subset of the training data points, which are called support vectors.
- *Support Vector Machine (SVM)*: is a decision machine, which is obtained by maximising the margins, thus it is also called maximum margin classifier.



- We have  $N$  training data vectors  $\mathbf{x}_1, \dots, \mathbf{x}_N$  and their target values  $t_1, \dots, t_N$  where  $t_N \in \{-1, 1\}$ .

- **SVM** for the two-class problem takes the form of

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b$$

where  $\phi(\mathbf{x})$  denotes a feature mapping,  $b$  the bias.

- Data points  $\mathbf{x}$  are classified by the sign of  $y(\mathbf{x})$ .
- Assume that the training data is linear separable in feature space. Thus, optimal  $\mathbf{w}$  and  $b$  satisfy

$$t_n y(\mathbf{x}_n) > 0$$


for all training data points.

- The perpendicular distance of a point  $\mathbf{x}$  from a hyperplane  $y(\mathbf{x})$  is  $|y(\mathbf{x})| / \|\mathbf{w}\|$ .
- As we assume  $t_n y(\mathbf{x}_n) > 0$  for all  $n$ , the distance of a point  $\mathbf{x}_n$  to the decision surface is

$$\frac{t_n y(\mathbf{x}_n)}{\|\mathbf{w}\|} = \frac{t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|}$$

- The margin is the minimum perpendicular distance.
- We find  $\mathbf{w}$  and  $b$  that maximise the margin i.e.

$$\arg \max_{\mathbf{w}, b} \left\{ \frac{1}{\|\mathbf{w}\|} \min_n [t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b)] \right\}$$



# Dual representation

- Canonical representation of the decision hyperplane:

*Rescaling  $\mathbf{w} \rightarrow k\mathbf{w}$  and  $b \rightarrow kb$  does not change the distance from any point  $\mathbf{x}_n$  to the decision surface. We can therefore set*

$$t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) = 1$$

for the point that is closest to the surface. All data points satisfy

$$t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) \geq 1, \quad n = 1, \dots, N.$$

# Dual representation

- The optimisation problem becomes to maximise  $||\mathbf{w}||^{-1}$ . Equivalently, we have

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} ||\mathbf{w}||^2$$

subject to the constraints  $t_n(\mathbf{w}^T \phi(\mathbf{x}) + b) \geq 1$

- The problem is solved by introducing Lagrange multipliers  $a_n \geq 0$ ,

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} ||\mathbf{w}||^2 - \sum_{n=1}^N a_n \{t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) - 1\}$$

where  $\mathbf{a} = (a_1, \dots, a_N)^T$ . Note the minus sign in front of the Lagrange multiplier term, as we are minimising the function.

$$\max f(x) \text{ s.t. } g(x)=0 \quad \longrightarrow \quad \max f(x) + \lambda g(x)$$



# Dual representation

- Setting the derivatives of  $L(\mathbf{w}, b, \mathbf{a})$  w.r.t.  $\mathbf{w}$ ,  $b$  to zero, we obtain

$$\mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \qquad 0 = \sum_{n=1}^N a_n t_n.$$

Using the two conditions above, we eliminate  $\mathbf{w}$  and  $b$  to get the *dual representation* as

$$\tilde{L}(\mathbf{a}) = \sum_{i=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

with the constraints  $a_n \geq 0, \quad n = 1, \dots, N$

$$\sum_{n=1}^N a_n t_n = 0.$$

where  $k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$ .

# Dual representation

- The SVM decision function takes the form of

$$y(\mathbf{x}) = \sum_{i=1}^N a_n t_n \boxed{k(\mathbf{x}, \mathbf{x}_n)} + b.$$

- *Karush-Kuhn-Tucker (KKT) conditions* of the Lagrangian function are

$$L(\mathbf{x}, \lambda) = f(\mathbf{x}) + \lambda g(\mathbf{x}). \quad \rightarrow \quad \begin{aligned} g(\mathbf{x}) &\geq 0, \\ \lambda &\geq 0, \\ \lambda g(\mathbf{x}) &= 0. \end{aligned}$$

- The KKT conditions of the maximum margin problem are

$$a_n \geq 0$$

$$t_n y(\mathbf{x}_n) - 1 \geq 0$$


$$a_n \{t_n y(\mathbf{x}_n) - 1\} = 0.$$



For every data point,  $a_n=0$  or  $t_n y(\mathbf{x}_n)=1$ .

Any data point for which  $a_n=0$  plays no role in making predictions.

The remaining data points are called *support vectors*, and because they satisfy  $t_n y(\mathbf{x}_n)=1$ , i.e. lying on the margin hyperplanes.




# Optimisation

- *Sequential minimal optimisation (SMO):*


see e.g.

CHRISTOPHER J.C. BURGESS, A Tutorial on Support Vector  
Machines for Pattern Recognition

(<http://research.microsoft.com/pubs/67119/svmtutorial.pdf>)



# Overlapping class distributions



# Overlapping class distributions

- Slack variables,  $\xi_n \geq 0$ ,  $n=1, \dots, N$  are defined s.t.

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n, \quad n = 1, \dots, N.$$

$\xi_n = 0$  for data points that are on or inside the correct margin boundary and  $\xi_n = |t_n - y(\mathbf{x}_n)|$  for other points.

- We therefore minimise

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

where  $C > 0$  is a trade-off constant and  $\xi_n \geq 0$ , subject to

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n, \quad n = 1, \dots, N.$$




# Probabilistic output

- We fit a logistic sigmoid to the output of SVM. The conditional probability takes the form of

$$p(t = 1|\mathbf{x}) = \sigma(Ay(\mathbf{x}) + B)$$

Where  $A$ ,  $B$  are parameters and  $\sigma$  is the sigmoid function.




# Multiclass SVMs 1)

**One-versus-the-rest** trains  $K$  separate SVMs, where  $y_k(\mathbf{x})$  is trained by  $C_k$  as the positive class and the remaining  $K-1$  classes as the negative class.

The prediction for a new input  $\mathbf{x}$  is by

$$y(\mathbf{x}) = \max_k y_k(\mathbf{x}).$$




Problems: 1) the output values  $y_k(\mathbf{x})$  for different classifiers have no appropriate scales. 2) the training sets are imbalanced.

# Multiclass SVMs 2)

**One-versus-one** trains  $K(K-1)/2$  different 2-class SVMs on all possible pairs of classes.

Classification of a test point is by, which class has the highest number of 'votes'.



*Problems: it requires more training time and evaluation time.*

# Statistical Pattern Recognition Toolbox for Matlab

[http://cmp.felk.cvut.cz/cmp/software/stp  
rtool/](http://cmp.felk.cvut.cz/cmp/software/stp<br/>rtool/)

# BoW as input to SVM classifier

- Treat BoW histograms as feature vectors for standard classifiers
  - e.g SVM
- SVM for object classification
  - Csurka, Bray, Dance & Fan, 2004



# Video-based Object Recognition



Imperial College, KAIST

# Video-based Object Recognition by Sets of Sets